

Victoria University of Wellington
School of Engineering and Computer Science

SWEN221: Software Development

Lab Handout (worth $\approx 1.3\%$ of overall mark)

The aim of this activity is to implement a *parallel quicksort* of integer data. This exploits multi-threading in order to do the quicksort computation more efficiently, particularly on machines with more than one processor (CPU).

A small library for controlling parallel code is provided in the `swen221.concurrent` package. This is a simpler version of the tools found in `java.util.concurrent`. You can download the `quicksort.jar` file from the SWEN221 homepage.

1 ParSum — Parallel Sum

To illustrate how the `swen221.concurrent` library works, a small example called `ParSum` is provided. This reads a file containing integers and sums them to produce a single result. You can run `ParSum` from the command-line like so:

```
% java ParSum large.dat
Executing on machine with 4 processor(s).
Read 10000000 data items.
Performing a SEQUENTIAL sum...
Summed data in 1200ms
Sum was 49999995000000
Data was summed correctly!
```

The above executes a standard sequential version of the sum operation (**NOTE:** you can execute this directly from Eclipse by providing command-line arguments to the run configuration). To execute the sum operation in parallel, you add the “`--parallel`” option:

```
% java ParSum --parallel large.dat
Executing on machine with 4 processor(s).
Read 10000000 data items.
Performing a PARALLEL sum...
Summed data in 870ms
Sum was 49999995000000
Data was summed correctly!
```

You should notice that the parallel version takes less time to execute on a machine with more than one processor. The amount of speed up will vary and (unfortunately) will not be linear in the number of processors. You will find the method `ParSum.parallelSum()` helpful in completing this lab.

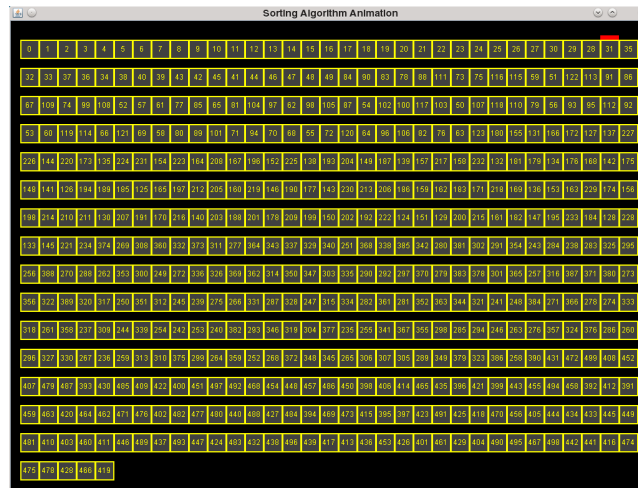


Figure 1: Illustrating the Sort GUI

2 ParSort — Parallel Quick Sort

An initial implementation of the quick sort code has been provided for you in the class ParSort. In particular, this contains the method `ParSort.sequentialSort()` which implements a *sequential* quick sort. You can run the ParSort code like so:

```
% java ParSort large.dat
Executing on machine with 4 processor(s).
Read 10000000 data items.
Performing a SEQUENTIAL quicksort...
Sorted data in 10756ms
Data was sorted correctly!
```

A special GUI is also provided which animates the quick sort. You can view this like so:

```
% java ParSort --gui tiny.dat
Executing on machine with 4 processor(s).
Read 500 data items.
Running in Display Mode (so ignore timings).
Performing a SEQUENTIAL quicksort...
```

You may find the GUI mode helpful in debugging your code (**NOTE:** you should use only the `tiny.dat` file, as others are too big). The GUI window is shown in Figure 1.

Finally, to execute the quick sort in parallel, you add the “`--parallel`” option:

```
% java ParSort --parallel large.dat
Executing on machine with 4 processor(s).
Read 10000000 data items.
Performing a PARALLEL quicksort...
Sorted data in 0ms
!!! ERROR: data not sorted correctly
```

NOTE: the result from the parallel sort is currently incorrect because the parallel implementation is not finished. Specifically, a method `ParSort.parallelSort()` is provided without any implementation and your task is to implement this!

2.1 HINTS

You should find the method `ParSum.parallelSum()` helpful in completing the implementation of `ParSort.parallelSort()`. However, there are some important differences in the problems and you will need to consider carefully how to parallelise the quick sort. There are two approaches you can take:

1. **Division based on Job Size.** This approach follows `ParSum.parallelSum()` quite closely. The input data list is split into a roughly even number of chunks, based on the number of processors available. A job is created for each chunk, and then the code waits for each Job to finished. *However, at this point, the results from each chunk must still be recombined to produce the final sorted list.*
2. **Division based on Depth.** This approach differs from `ParSum.parallelSum()`, but is easier to implement. In this case, `sequentialSort()` is modified to count the depth of the computation tree. Initially, the computation proceeds sequentially as before. But, at a certain depth (e.g. 2), a Job is created to complete the remainder of the chunk being sorted. This way, you end up creating 2^n Jobs at a depth n . *The advantage of this approach is that when the Jobs are finished, the data is sorted correctly.*

When your implementation of `ParSort.parallelSort()` is working correctly, you should find that it outperforms `ParSort.sequentialSort()` on large data sets.

Marking Guide

Each lab is worth just under 1% of your overall mark for SWEN221. The lab should be marked during the lab sessions, according to the following grade scale:

- **0:** Student didn't attend lab.
- **E:** Student did not really participate in the lab.
- **D:** Student's participation was *poor*. For example, he/she made some attempt to work on the lab, but did not complete any activities.
- **C:** Student's participation was *satisfactory*. That is, he/she attempted the lab but didn't achieve any parallelisation.
- **B:** Student's participation was *good*. That is, he/she got parallelisation working, but result produced was incorrect.
- **A:** Student's participation was *excellent*. That is, he/she computed the correct answer, and it was faster than sequential version on `large.dat`.